

Designing Reservation Reform Without Mass Unrest

An agent-based simulation of protest mobilization against an income-only quota in India, and of the policy designs that would keep it small

Utsav Balhara

B.Tech, Netaji Subhas University of Technology (NSUT), New Delhi

September 2026 · github.com/utsavbalhara/reservation-reform-protest-model

At a glance

Question. If India replaced caste-based reservation with a single income test (family income below ₹8 lakh), how many people would protest, and which policy designs would keep that number small?

Method. A synthetic India of 120,000 agents (each standing for about 12,000 people) whose grievance combines loss-averse material loss with symbolic threat to caste recognition. Protest spreads through neighbourhood and national threshold cascades amplified by party mobilization on bandh days. Thirteen scenarios were each run 50 times with paired uncertainty draws, in two mobilization regimes.

Findings.

- An abrupt switch draws about 46.1 lakh onto the streets on the peak day (1.29 crore ever protest, about 18.5 deaths).
- Keeping caste as the qualifier and adding an income filter inside each quota cuts peak turnout by −94%. Most of the anger is about losing caste recognition, not about losing a seat.
- If the income-only end state is kept, a data-first commission with cross-party consensus is the strongest single lever (−82%). A managed-transition package reaches 2.2 lakh (−95%).
- Compensating the people who lose eligibility does little (−24%).
- Internet shutdowns and heavy policing lower the peak modestly while multiplying deaths by 2.9 and 4.3 times.

Status of the numbers. Magnitudes depend on assumed parameters; the ranking of interventions is the robust result.

Contents

1	The question	2
1.1	Who actually changes status	2
2	Theory: why people would protest	2
2.1	Loss aversion and the politics of taking things away	2
2.2	Symbolic threat and group identity	3
2.3	Threshold cascades	3
2.4	Organization and elite cues	3
2.5	Repression and backfire	3
3	The model in simple terms	4
3.1	The synthetic India	4
3.2	How one day works	4
3.3	Calibration and a validation check	5
4	Results at a glance	6
5	Scenario by scenario	7
5.1	Base: the abrupt ₹8 lakh income-only switch	7
5.2	V: validation, symbolic threat only	8
5.3	L1: grandfather current cohorts, with a 10-year glide	9
5.4	L2: expand seats so no group loses	9
5.5	L3: keep caste sub-quotas, add an income filter inside them	10
5.6	L4: sub-classify to favour the most-deprived	10
5.7	L5: data-first commission with cross-party consensus	11
5.8	L6: compensate above-line losers	11
5.9	L7: guarantee untouched protections	12
5.10	B1: internet shutdowns (tested for backfire)	12
5.11	B2: heavy policing and mass arrests (tested for backfire)	13
5.12	C1: the managed-transition package	14
5.13	C2: the hybrid-design package	14
6	Robustness	15
7	What this means	15
7.1	Limitations	15
8	Reproducing this report	16

1 The question

India reserves seats in public education and government jobs for Scheduled Castes (SC, 15%), Scheduled Tribes (ST, 7.5%) and Other Backward Classes (OBC, 27%). Since 2019 it has also reserved 10% for Economically Weaker Sections (EWS) of the general category. Two of these categories already carry an income test. OBC candidates from families above ₹8 lakh a year (the “creamy layer”) are excluded, and EWS applies only below ₹8 lakh. SC and ST quotas have no income test at all.

This report asks what would happen if the whole system were converted into a single income test at ₹8 lakh, with caste removed as a criterion. It asks two questions:

1. How many people would protest, and for how long?
2. Which changes to the policy’s design, sequencing and politics would keep that protest small?

1.1 Who actually changes status

At the ₹8 lakh line the eligibility change is small. OBC families above the line and general-category families below it already have the status an income test would give them. The people who lose eligibility are SC and ST families above ₹8 lakh, about 2.8% of the population in this model (roughly 4 crore people). About 80% of Indians are reserved before and after, and about 17% are unreserved before and after.

That small eligibility shift is why protest is a puzzle worth modelling. Recent history shows that reservation politics mobilizes around what a change symbolizes, not only around who loses a benefit:

- In 2018 a Supreme Court ruling that diluted the SC/ST (Prevention of Atrocities) Act removed nobody’s quota, yet it triggered the largest Bharat Bandh on record and at least 11 deaths.
- In 2024 the Supreme Court permitted sub-classification within SC/ST quotas, and some judges suggested a creamy layer for them. That triggered another nationwide bandh.

2 Theory: why people would protest

Five bodies of theory shape the model. Each is stated formally, then in plain words.

2.1 Loss aversion and the politics of taking things away

People weigh losses roughly twice as heavily as equal gains; Tversky and Kahneman estimate a loss-aversion coefficient of about 2.25. Pierson’s work on welfare retrenchment builds on this. Cutting an existing benefit is politically harder than creating one, because the losses are concentrated, immediate and visible while the gains are diffuse and uncertain. Governments that succeed at retrenchment typically use three techniques:

- *Grandfathering*, which exempts people already receiving the benefit.
- *Compensation* for the people who lose.
- *Blame-sharing* through broad coalitions.

In plain words. Taking away a quota feels worse than never having had one, and the people who lose are easier to organize than the people who gain.

2.2 Symbolic threat and group identity

Social identity theory holds that people derive part of their self-worth from group membership, and react strongly to threats to the group's status. For SC and ST communities, caste-based reservation is also a constitutional recognition of historical injustice. Removing caste as the basis therefore carries a symbolic threat independent of any individual's eligibility. The 2018 bandh is the clearest Indian evidence that symbolic threat alone can mobilize at national scale.

In plain words. Many people would protest not because they lose a seat, but because they read the change as erasing their community's recognition.

2.3 Threshold cascades

Granovetter's threshold model treats each person as having a personal threshold: the share of others who must already be protesting before they join. A few activists join at once, and their presence tips people with slightly higher thresholds, and so on. Small differences in the distribution of thresholds or grievances can produce very large differences in turnout, which makes the system non-linear near its tipping point.

In plain words. People join when enough of their neighbours already have. Pushing grievance down a little can remove far more protest than you would expect, because it stops the chain reaction early.

2.4 Organization and elite cues

Resource-mobilization theory (McCarthy and Zald) argues that grievance is common but organized protest is not. Organizations, parties and networks supply the buses, bandh calls and messaging that convert grievance into turnout. Both the 2018 and 2024 bandhs had support from opposition parties.

In plain words. Anger needs machinery. When opposition parties co-own a reform, they stop supplying that machinery to protests against it.

2.5 Repression and backfire

Studies of repression find that visible force against protesters often backfires. Deaths create martyrs, and perceived injustice broadens support. India has imposed more internet shutdowns than any other country, and Rydzak finds that shutdowns in India are associated with more violent collective action, not less, because they disrupt the coordination that peaceful protest relies on.

In plain words. Cutting the internet or sending in heavy policing may shrink the crowd on a given day, but it tends to make protest angrier, longer and deadlier.

3 The model in simple terms




Figure 1: What drives one person's decision each day. Dashed orange arrows are feedback loops: today's turnout shapes tomorrow's decisions.

3.1 The synthetic India

The model builds 120,000 agents, each standing for about 12,000 people. Group shares follow the 2011 Census for SC (16.6%) and ST (8.6%) and survey estimates for OBC (42%), leaving 32.8% General. About one family in five sits above ₹8 lakh: 30% of General, 17% of OBC, 12% of SC and 10% of ST. SC and ST are split into a better-off tier and a most-deprived tier, with most above-line families in the better-off tier. Agents live in neighbourhoods of about 100 agents from the same group.

3.2 How one day works

For every agent, every day:

1. **Grievance** is the material loss (weighted 2.25 times if it is a loss) plus the symbolic threat multiplied by how strongly the agent identifies with its group.
2. **Social pull** rises with the share of the agent's neighbourhood that protested yesterday and with its group's national turnout. Both effects saturate, because a crowd can only pull so hard.
3. **Organization** adds a push on bandh-call days (days 3, 13 and 26), larger when opposition parties back the protest.
4. The agent **protests** with a probability that rises smoothly as grievance plus pull plus push exceeds its personal threshold.
5. Protesting raises the agent's threshold a little (**fatigue**).
6. Deaths occur in proportion to turnout. Each death raises symbolic threat for the groups on the street (the **martyr effect**).

Formally, agent i in group g protests on day t with probability

$$P_{i,t} = \sigma \left(\frac{w_m L(m_i) + w_s \iota_i(s_g + \mu_{g,t}) + \beta_n \tanh \frac{n_{c(i),t-1}}{n^*} + \beta_v \tanh \frac{v_{g,t-1}}{v^*} + a_{og} b_t - \theta_{i,t}}{\tau} \right), \quad (1)$$

where $L(m) = \lambda m$ for losses and m for gains, ι_i is identity strength, s_g symbolic threat, $\mu_{g,t}$ the martyr effect, n and v neighbourhood and national turnout, o_g organizational capacity, a the opposition-party amplifier, b_t the bandh-day mobilization, and $\theta_{i,t}$ the agent's threshold.

3.3 Calibration and a validation check

The one free level parameter is the mean participation threshold. It is set to 6.6, which puts the baseline peak near 46 lakh, the middle of an independent prior estimate of 20 lakh to 1 crore. The check that matters is separate. Removing all material loss leaves only symbolic threat, much like the 2018 Atrocities Act ruling. That shock alone should still produce a bandh at 2018 scale. The model gives 13.3 lakh, which passes.




Figure 2: Peak-day turnout against the mean threshold. Turnout roughly halves for every quarter-step in threshold. Any lever that lowers net motivation therefore has a multiplied effect on turnout.




Figure 3: Average grievance by segment in the baseline. Symbolic threat is the larger part for every group except SC and ST families above the line, who make up under 3% of India.

4 Results at a glance





Figure 4: Peak-day protesters for all thirteen scenarios, median of 50 paired runs with the 10th–90th percentile range, on a log scale.

Table 1: Central regime, 50 paired runs per scenario. Peak in lakh; cumulative in crore.

	Scenario	Peak	10th–90th	Change	Ever protest	Deaths
Base	Abrupt ₹8L income-only switch	46.1	14.9–118.9	–	1.29	18.5
V	Validation: symbolic threat only (2018-type shock)	13.3	6.0–37.6	–71%	0.47	5.5
L1	Grandfather current cohorts + 10-year glide	15.6	6.6–41.5	–66%	0.50	5.5
L2	Expand seats so no group loses	24.1	9.8–65.8	–48%	0.74	11
L3	Keep caste sub-quotas, add income filter	2.9	1.7–4.3	–94%	0.10	1
L4	Sub-classify to favour most-deprived	26.9	11.6–71.4	–42%	0.83	13.5
L5	Data-first commission + cross-party consensus	8.3	4.1–24.2	–82%	0.41	4.5
L6	Compensate above-line losers	34.9	13.1–93.9	–24%	1.08	16.5
L7	Guarantee untouched protections	18.2	9.6–55.2	–61%	0.68	7.5
B1	Internet shutdowns	42.8	17.2–73.7	–7%	1.41	54.5
B2	Heavy policing and mass arrests	32.9	17.4–73.3	–29%	1.61	79.5
C1	Managed transition package (L1+L2+L5+L6+L7)	2.2	1.4–3.4	–95%	0.12	1
C2	Hybrid design package (L3+L4+L1+L5+L6+L7)	0.6	0.4–0.9	–99%	0.04	0

Table 2: High-mobilization regime (mean threshold lowered by 0.5), 50 paired runs per scenario.

	Scenario	Peak	10th–90th	Change	Ever protest	Deaths
Base	Abrupt ₹8L income-only switch	209.4	92.6–388.4	–	4.14	79
V	Validation: symbolic threat only (2018-type shock)	66.1	18.0–152.5	–68%	1.84	27
L1	Grandfather current cohorts + 10-year glide	74.2	20.0–160.7	–65%	1.98	30.5
L2	Expand seats so no group loses	117.0	35.5–258.0	–44%	2.50	45
L3	Keep caste sub-quotas, add income filter	7.7	5.1–12.5	–96%	0.29	3
L4	Sub-classify to favour most-deprived	106.5	40.0–225.0	–49%	2.54	46.5
L5	Data-first commission + cross-party consensus	42.2	13.5–106.7	–80%	1.55	26.5
L6	Compensate above-line losers	167.0	49.3–327.2	–20%	3.80	65
L7	Guarantee untouched protections	110.8	34.9–247.2	–47%	2.52	43
B1	Internet shutdowns	110.6	55.9–184.4	–47%	3.43	152.5
B2	Heavy policing and mass arrests	124.3	52.4–245.6	–41%	4.20	254.5
C1	Managed transition package (L1+L2+L5+L6+L7)	6.4	3.6–10.0	–97%	0.36	4
C2	Hybrid design package (L3+L4+L1+L5+L6+L7)	1.9	1.3–2.9	–99%	0.12	1

**Figure 5:** Daily turnout for four contrasting scenarios. Bandh calls produce the spikes. Heavy policing flattens those spikes but keeps more people on the street between them.

5 Scenario by scenario

Each scenario below follows the same layout: what the policy is, the theory behind it, a plain-language summary, how it enters the model, the outcome, and why the outcome came out that way. All percentages compare peak-day turnout with the baseline median of 46.1 lakh.

5.1 Base: the abrupt ₹8 lakh income-only switch

What it is

Parliament replaces SC, ST, OBC and EWS quotas with a single income test at ₹8 lakh, effective immediately, with no consultation or transition.

Theory

Every protest driver is at full strength. Above-line SC and ST families suffer a concentrated loss; below-line reserved groups lose their separate pools; symbolic threat is unmitigated; and opposition parties are free to mobilize.

In plain words. This is the reference point: the reform done in the most provocative way.

Outcome

Peak day (central)	Change	Ever protested	Deaths
46.1 lakh (14.9–118.9)	–	1.29 crore	18.5

High-mobilization regime: 2.09 crore peak, 4.14 crore ever protest, 79 deaths.

Why

SC communities carry most of the turnout (Figure 6): they have the highest symbolic threat, the strongest organizations and the largest group of direct losers. Turnout spikes on each bandh call and decays between them as fatigue sets in. The wide range across runs reflects how close the system sits to its tipping point.




Figure 6: Baseline turnout by group. SC turnout dominates; General-category turnout is zero because the reform favours them.

5.2 V: validation, symbolic threat only

What it is

A check, not a policy. All material loss is removed, so only the symbolic threat of abolishing caste-based quotas remains. This resembles the 2018 Atrocities Act ruling, which changed no one's quota.

Theory

If identity-based mobilization is real, a symbolic shock alone should still produce a large national bandh.

In plain words. We checked that the model can reproduce something like the 2018 bandh before trusting it on anything else.

Outcome

Peak day (central)	Change	Ever protested	Deaths
13.3 lakh (6.0–37.6)	–71%	0.47 crore	5.5

High-mobilization regime: 66.1 lakh peak.

Why

Symbolic threat alone produces bandh-scale protest, consistent with 2018. Removing material loss still cuts turnout by about 70%, because material loss pushes the most motivated agents over their thresholds and they seed the cascade.

5.3 L1: grandfather current cohorts, with a 10-year glide

What it is

Students and job applicants already in the pipeline keep their current category rights. The income test is phased in over ten years.

Theory

Loss aversion applies to losses felt now. Grandfathering is Pierson’s classic blame-avoidance tool: it converts an immediate, concentrated loss into a deferred and discounted one.

In plain words. Nobody who is counting on a quota today loses it, and the change arrives slowly enough to plan around.

How it is modelled

Every material loss is cut to 30% of its baseline value; symbolic threat falls by 5%.

Outcome

Peak day (central)	Change	Ever protested	Deaths
15.6 lakh (6.6–41.5)	–66%	0.50 crore	5.5

High-mobilization regime: 74.2 lakh (–65%).

Why

Grandfathering takes out most of the material component, so turnout falls almost to the symbolic-only level (13.3 lakh). What remains is protest about recognition, which grandfathering does not address.

5.4 L2: expand seats so no group loses

What it is

Seats and posts are expanded so that no group’s absolute number of seats falls, as the government did in 2019 when it added about 2.14 lakh seats in central institutions to introduce EWS without cutting other quotas.

Theory

When reserved pools merge, below-line SC, ST and OBC candidates face more competition. Expanding capacity removes that loss for the large below-line majority.

In plain words. If the pie grows, nobody’s slice has to shrink.

How it is modelled

Material loss for below-line SC, ST and OBC is set to zero. Above-line SC and ST still lose eligibility.

Outcome

Peak day (central)	Change	Ever protested	Deaths
24.1 lakh (9.8–65.8)	–48%	0.74 crore	11

High-mobilization regime: 1.17 crore (–44%).

Why

The below-line loss is spread thinly across many people, so removing it lowers grievance only modestly per person. Both the direct losers and the symbolic threat remain. It is a useful supporting measure, not a solution.

5.5 L3: keep caste sub-quotas, add an income filter inside them

What it is

A hybrid design. SC, ST and OBC quotas stay, but within each quota only families below the income line are eligible. In effect this extends the OBC creamy layer to SC and ST. Brazil's 2012 Quota Law works similarly: racial sub-quotas sit inside a public-school and income filter.

Theory

It keeps caste recognition, which carries most of the symbolic threat, while still targeting benefits by income. It also avoids merging pools, so below-line candidates lose nothing.

In plain words. Keep the caste categories, but stop the richest families inside them from using the quota.

How it is modelled

Symbolic threat for SC, ST and OBC falls to 40% of baseline; below-line material loss is removed. Above-line SC and ST still lose eligibility.

Outcome

Peak day (central)	Change	Ever protested	Deaths
2.9 lakh (1.7–4.3)	–94%	0.10 crore	1

High-mobilization regime: 7.7 lakh (–96%).

Why

It attacks the largest driver directly, and the cascade does the rest: once most agents fall below threshold, the chain reaction never starts. Remaining protest comes from above-line SC and ST families, who still lose. The result holds even if the hybrid removes only a fifth of the symbolic threat (9.2 lakh; Section 6). The main caution is the 2024 bandh, which a mere suggestion of an SC/ST creamy layer triggered. The real symbolic relief may be at the weak end of the tested range. This design also gives up the premise of the question, since caste stays a criterion.

5.6 L4: sub-classify to favour the most-deprived

What it is

Within SC and ST, a larger share of benefits goes to the most-deprived sub-castes, as permitted by the Supreme Court in 2024.

Theory

Coalitions can split when a reform creates real winners inside a protesting group. During the 2024 bandh, about 30 SC/ST communities in Rajasthan opposed the protest because they supported sub-classification.

In plain words. Give the poorest sub-castes a better deal, and many of them will not join the protest.

How it is modelled

The most-deprived tier gains (material loss reduced by 0.5) and feels only 30% of the symbolic threat.

Outcome

Peak day (central)	Change	Ever protested	Deaths
26.9 lakh (11.6–71.4)	–42%	0.83 crore	13.5

High-mobilization regime: 1.06 crore (–49%).

Why

The most-deprived tier largely drops out, but the better-off tier, which includes nearly all the direct losers, keeps protesting. Many neighbourhoods are mixed across tiers, so the better-off tier's cascade still runs.

5.7 L5: data-first commission with cross-party consensus

What it is

The reform goes through a statutory commission that works from the caste census data before any legislation, with opposition parties represented and the final design negotiated across parties.

Theory

It combines procedural legitimacy with Pierson's blame-sharing. When opposition parties co-own a reform, they do not supply the organization, messaging and bandh calls that turn grievance into turnout.

In plain words. If every major party signs off, nobody has a reason to bus people into the streets against it.

How it is modelled

The opposition-party amplifier falls from 1.5 to 1.0; symbolic threat falls by 20%.

Outcome

Peak day (central)	Change	Ever protested	Deaths
8.3 lakh (4.1–24.2)	–82%	0.41 crore	4.5

High-mobilization regime: 42.2 lakh (–80%).

Why

Bandh-day mobilization sets peak turnout, and consensus removes a third of it while lowering symbolic threat for everyone. This is the strongest lever that keeps the full income-only end state, and it is the load-bearing part of the managed-transition package. It is also the hardest to achieve in practice.

5.8 L6: compensate above-line losers

What it is

Families that lose eligibility receive scholarships, hostel places and transitional grants.

Theory

Pierson's compensation strategy: buy off the concentrated losers.

In plain words. Pay the people who lose the quota.

How it is modelled

Material loss for above-line SC and ST is halved.

Outcome

Peak day (central)	Change	Ever protested	Deaths
34.9 lakh (13.1–93.9)	–24%	1.08 crore	16.5

High-mobilization regime: 1.67 crore (–20%).

Why

The direct losers are under 3% of the population, and money does nothing about symbolic threat. Compensation is the weakest lever, and removing it from the managed-transition package barely changes that package's result.

5.9 L7: guarantee untouched protections

What it is

The reform is written with explicit guarantees about what stays untouched: SC/ST reserved seats in legislatures, the Atrocities Act, and a scheduled review clause.

Theory

Symbolic threat depends on what people believe the reform will lead to next. Credible guarantees limit that fear.

In plain words. Say clearly, in law, what is not changing.

How it is modelled

Symbolic threat falls by 15% for every group.

Outcome

Peak day (central)	Change	Ever protested	Deaths
18.2 lakh (9.6–55.2)	–61%	0.68 crore	7.5

High-mobilization regime: 1.11 crore (–47%).

Why

A 15% cut in symbolic threat produces a much larger cut in turnout. Symbolic threat is the dominant driver, and the cascade multiplies any reduction.

5.10 B1: internet shutdowns (tested for backfire)

What it is

Mobile internet is suspended in areas with protests, as Indian states have frequently done.

Theory

Rydzak finds that shutdowns in India are associated with more violent collective action, because they disrupt the coordination that disciplined, peaceful protest needs.

In plain words. Cutting the internet makes it harder to organize, but it makes what does happen more chaotic and violent.

How it is modelled

Neighbourhood and national influence fall by 40% and bandh mobilization by 20%, while the death rate rises 2.5 times.

Outcome

Peak day (central)	Change	Ever protested	Deaths
42.8 lakh (17.2–73.7)	–7%	1.41 crore (+9%)	54.5 (×2.9)
High-mobilization regime: 1.11 crore (–47%), 152.5 deaths (×1.9).			

Why

Weaker coordination trims the peak, and more so when mobilization is high. But extra deaths feed the martyr effect, which raises symbolic threat and keeps people coming out, so more people protest overall in the central regime. Every version of this lever costs lives.

5.11 B2: heavy policing and mass arrests (tested for backfire)

What it is

Large deployments, preventive detentions and forceful dispersal on bandh days.

Theory

Force raises the cost of turning out but produces martyrs and perceived injustice, which broadens support.

In plain words. Force can shrink a crowd on the day, but every death recruits more protesters for the days after.

How it is modelled

Turning out on bandh days costs 0.4 more; the death rate triples; each death raises symbolic threat 1.5 times as much.

Outcome

Peak day (central)	Change	Ever protested	Deaths
32.9 lakh (17.4–73.3)	–29%	1.61 crore (+25%)	79.5 (×4.3)
High-mobilization regime: 1.24 crore (–41%), 254.5 deaths (×3.2).			

Why

Figure 5 shows the mechanism. Policing flattens the bandh-day spikes, but the martyr effect keeps more people on the street between bandhs, so more people protest in total and many more die.




Figure 7: Suppression indexed to the baseline (= 100), central regime.

5.12 C1: the managed-transition package

What it is

Keep the full income-only end state, but get there with grandfathering (L1), seat expansion (L2), a consensus commission (L5), compensation (L6) and guarantees (L7).

Theory

Each lever pushes agents further below threshold. Because turnout responds non-linearly, combined levers do better than their separate effects suggest.

In plain words. Do the same reform, but slowly, with more seats, with every party on board, and with clear promises.

Outcome

Peak day (central)	Change	Ever protested	Deaths
2.2 lakh (1.4–3.4)	–95%	0.12 crore	1

High-mobilization regime: 6.4 lakh (–97%).

Why

The package stays below the tipping point in both regimes. Leaving out the consensus commission triples turnout (6.6 lakh against 2.2 lakh), while leaving out compensation barely matters (2.4 lakh).

5.13 C2: the hybrid-design package

What it is

The hybrid design (L3) with sub-classification (L4), grandfathering (L1), the consensus commission (L5), compensation (L6) and guarantees (L7).

In plain words. Keep caste categories with an income filter, favour the poorest sub-castes, phase it in, and get every party on board.

Outcome

Peak day (central)	Change	Ever protested	Deaths
0.6 lakh (0.4–0.9)	–99%	0.04 crore	0

High-mobilization regime: 1.9 lakh (–99%).

Why

Almost every driver is addressed at once. What remains is roughly the protest an ordinary contested policy change would draw.

6 Robustness

Two checks test whether the headline findings rest on a single assumption.

- **Hybrid sensitivity.** The hybrid was rerun leaving 40%, 60% and 80% of symbolic threat in place. Peak turnout was 2.7 lakh, 4.6 lakh and 9.2 lakh: still at least an 80% cut.
- **Managed-transition leave-one-out.** Each of the five components was removed in turn. Only removing the consensus commission changes the result materially.

The ranking of levers is the same in the central and high-mobilization regimes.




Figure 8: Left: hybrid sensitivity. Right: leave-one-out on the managed-transition package (30 paired runs each).

7 What this means

1. **Design matters more than compensation.** Protest against this reform is mostly about recognition, so a design that preserves caste recognition while targeting by income does far better than paying the people who lose.
2. **Process is policy.** If the income-only end state is the goal, a data-first commission with cross-party ownership matters more than any single material concession.
3. **Suppression is not a mitigation strategy.** In this model it lowers peak turnout only modestly while multiplying deaths, consistent with India's own record on shutdowns.
4. **Levers compound.** Because turnout responds non-linearly, packages do far better than their parts.

7.1 Limitations

- Grievance weights, lever effects and neighbourhood structure are assumptions informed by qualitative evidence, not estimates from Indian protest data.

- The calibration target is itself an estimate.
- Regional variation, counter-protests, courts, elections and media are not modelled.
- Neighbourhoods are single-group. Real mixing across groups could either dampen or spread protest.
- The model runs 40 days and says nothing about long-term political consequences.

Treat magnitudes as order-of-magnitude and the ranking as the finding.

8 Reproducing this report

```
pip install -r requirements.txt
python -m experiments.compare_interventions --regime central
python -m experiments.compare_interventions --regime high-mobilization
python -m experiments.check_robustness
python -m experiments.calibrate_participation_threshold
python -m experiments.record_daily_trajectories
python -m visualization.render_result_figures
python -m visualization.render_campaign_videos
python -m visualization.write_latex_result_macros
cd report && xelatex research_report.tex && xelatex research_report.tex
```

References

- Granovetter, M. (1978). Threshold models of collective behavior. *American Journal of Sociology*, 83(6), 1420–1443.
- Hess, D., & Martin, B. (2006). Repression, backfire, and the theory of transformative events. *Mobilization*, 11(2), 249–267.
- Kahneman, D., & Tversky, A. (1979). Prospect theory: An analysis of decision under risk. *Econometrica*, 47(2), 263–291.
- McCarthy, J. D., & Zald, M. N. (1977). Resource mobilization and social movements: A partial theory. *American Journal of Sociology*, 82(6), 1212–1241.
- Pierson, P. (1994). *Dismantling the Welfare State? Reagan, Thatcher, and the Politics of Retrenchment*. Cambridge University Press.
- Pierson, P. (1996). The new politics of the welfare state. *World Politics*, 48(2), 143–179.
- Rydzak, J. (2019). Of blackouts and bandhs: The strategy and structure of disconnected protest in India. SSRN Working Paper 3330413.
- Tajfel, H., & Turner, J. C. (1979). An integrative theory of intergroup conflict. In W. G. Austin & S. Worchel (Eds.), *The Social Psychology of Intergroup Relations*. Brooks/Cole.
- Tversky, A., & Kahneman, D. (1992). Advances in prospect theory: Cumulative representation of uncertainty. *Journal of Risk and Uncertainty*, 5(4), 297–323.
- Government of India, affidavit in *Janhit Abhiyan v. Union of India* (2022) on 2.14 lakh additional seats for EWS, as reported by The Tribune and Bar & Bench.
- Brazil, Law No. 12.711 of 29 August 2012 (Quota Law).
- State of Punjab v. Davinder Singh* (2024), Supreme Court of India, on sub-classification of Scheduled Castes.